

# DESCRIPTION

Reduce the time a Mercedes-Benz spends on the test bench

## Problem Statement Scenario:

Since the first automobile, the Benz Patent Motor Car in 1886, Mercedes-Benz has stood for important automotive innovations. These include the passenger safety cell with a crumple zone, the airbag, and intelligent assistance systems. Mercedes-Benz applies for nearly 2000 patents per year, making the brand the European leader among premium carmakers. Mercedes-Benz is the leader in the premium car industry. With a huge selection of features and options, customers can choose the customized Mercedes-Benz of their dreams.

To ensure the safety and reliability of every unique car configuration before they hit the road, the company's engineers have developed a robust testing system. As one of the world's biggest manufacturers of premium cars, safety and efficiency are paramount on Mercedes-Benz's production lines. However, optimizing the speed of their testing system for many possible feature combinations is complex and time-consuming without a powerful algorithmic approach.

You are required to reduce the time that cars spend on the test bench. Others will work with a dataset representing different permutations of features in a Mercedes-Benz car to predict the time it takes to pass testing. Optimal algorithms will contribute to faster testing, resulting in lower carbon dioxide emissions without reducing Mercedes-Benz's standards.

```
In [1]: #Importing Libraries

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: #Importing datasets
train_data=pd.read_csv('train1.csv')
test_data=pd.read_csv('test1.csv')
```

```
In [3]: #Checking shape of the datasets
train_data.shape, test_data.shape
```

```
Out[3]: ((4209, 378), (4209, 377))
```

```
In [4]: #Checking datatypes
train_data.dtypes
```

```
Out[4]: ID          int64
y          float64
X0          object
X1          object
X2          object
...
X380        int64
X382        int64
```

```
X383      int64
X384      int64
X385      int64
Length: 378, dtype: object
```

```
In [5]: # Printing columns of object datatypes
for i in train_data.columns:
    data_type = train_data[i].dtype
    if data_type == 'object':
        print(i)
```

```
X0
X1
X2
X3
X4
X5
X6
X8
```

**If for any column(s), the variance is equal to zero, then you need to remove those variable(s)**

```
In [6]: variance = pow(train_data.drop(columns={'ID', 'y'}).std(),2).to_dict()

null_cnt = 0
for key, value in variance.items():
    if(value==0):
        print('Name = ',key)
        null_cnt = null_cnt+1
print('No of columns which has zero variance = ',null_cnt)
```

```
Name = X11
Name = X93
Name = X107
Name = X233
Name = X235
Name = X268
Name = X289
Name = X290
Name = X293
Name = X297
Name = X330
Name = X347
No of columns which has zero variance = 12
```

Since these 12 columns have 0 variance, these all are need to be dropped

```
In [7]: #Dropping 12 columns with 0 variance from train dataset
train_data.drop(['X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289', 'X290', 'X293', 'X330', 'X297', 'X347'])
```

```
In [8]: train_data.shape
```

```
Out[8]: (4209, 366)
```

```
In [9]: #Dropping 12 columns with 0 variance from test dataset
test_data.drop(['X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289', 'X290', 'X293', 'X330', 'X297', 'X347'])
```

```
In [10]: test_data.shape
```

```
Out[10]: (4209, 365)
```

# Check for null and unique values for test and train sets

```
In [11]: #Null in train dataset
train_data.isnull().sum().any()
```

```
Out[11]: False
```

```
In [12]: #checking null in test dataset
test_data.isnull().sum().any()
```

```
Out[12]: False
```

It seems our datasets dont have any null values

```
In [13]: #Cheching Unique values in train dataset
train_data['X0'].unique
```

```
Out[13]: <bound method Series.unique of 0          k
1          k
2          az
3          az
4          az
..
4204       ak
4205       j
4206       ak
4207       al
4208       z
Name: X0, Length: 4209, dtype: object>
```

```
In [14]: #Cheching Unique values in test dataset
test_data['X0'].unique
```

```
Out[14]: <bound method Series.unique of 0          az
1          t
2          az
3          az
4          w
..
4204       aj
4205       t
4206       y
4207       ak
4208       t
Name: X0, Length: 4209, dtype: object>
```

```
In [15]: train_data['X1'].unique
```

```
Out[15]: <bound method Series.unique of 0          v
1          t
2          w
3          t
4          v
..
4204       s
4205       o
4206       v
4207       r
4208       r
Name: X1, Length: 4209, dtype: object>
```

```
In [16]: test_data['X1'].unique
```

```
Out[16]: <bound method Series.unique of 0          v
```

```

1      b
2      v
3      l
4      s
..
4204   h
4205   aa
4206   v
4207   v
4208   aa
Name: X1, Length: 4209, dtype: object>

```

## Apply label encoder

```

In [17]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()

```

```

In [18]: #Assigning feature and target in train dataset
train_data_feature = train_data.drop(['y', 'ID'], axis=1)
train_data_target = train_data['y']
print(train_data_feature.shape)
print(train_data_target.shape)

(4209, 364)
(4209,)

```

```

In [19]: #Describing object type variables in train dataset
train_data_feature.describe(include='object')

```

```

Out[19]:

```

	X0	X1	X2	X3	X4	X5	X6	X8
<b>count</b>	4209	4209	4209	4209	4209	4209	4209	4209
<b>unique</b>	47	27	44	7	4	29	12	25
<b>top</b>	z	aa	as	c	d	w	g	j
<b>freq</b>	360	833	1659	1942	4205	231	1042	277

```

In [20]: train_data_feature['X0'] = le.fit_transform(train_data_feature.X0)
train_data_feature['X1'] = le.fit_transform(train_data_feature.X1)
train_data_feature['X2'] = le.fit_transform(train_data_feature.X2)
train_data_feature['X3'] = le.fit_transform(train_data_feature.X3)
train_data_feature['X4'] = le.fit_transform(train_data_feature.X4)
train_data_feature['X5'] = le.fit_transform(train_data_feature.X5)
train_data_feature['X6'] = le.fit_transform(train_data_feature.X6)
train_data_feature['X8'] = le.fit_transform(train_data_feature.X8)

```

## Perform dimensionality reduction

```

In [21]: from sklearn.decomposition import PCA
pca = PCA(n_components=.95)

```

```

In [22]: #For Train dataset
pca.fit(train_data_feature, train_data_target)

```

```

Out[22]: PCA(n_components=0.95)

```

```

In [23]: train_data_feature_tf = pca.fit_transform(train_data_feature)
print(train_data_feature_tf.shape)

(4209, 6)

```

# Predict your test\_df values using XGBoost

## Building model using the train data set.

```
In [24]: import xgboost as xgb
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score, mean_squared_error
from math import sqrt
```

```
In [25]: train_x, test_x, train_y, test_y = train_test_split(train_data_feature_tf, train_data_target)
print(train_x.shape)
print(train_y.shape)
print(test_x.shape)
print(test_y.shape)
```

(2946, 6)

(2946,)

(1263, 6)

(1263,)

## XGBoost's hyperparameters tuning manually

```
In [26]: xgb_reg = xgb.XGBRegressor(objective='reg:linear', colsample_bytree = 0.3, learning_rate=0.1,
                                   n_estimators = 20)
model = xgb_reg.fit(train_x, train_y)
print('RMSE = ', sqrt(mean_squared_error(model.predict(test_x), test_y)))
```

[17:02:32] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-030221e36e1a46bfb-1/xgboost/xgboost-ci-windows/src/objective/regression\_obj.cu:213: reg:linear is now deprecated in favor of reg:squarederror.

RMSE = 12.288794806074309


RMSE = 12.288794806074309

```
In [27]: pred_test_y = model.predict(test_x)

plt.figure(figsize=(10,5))

sns.distplot(test_y[test_y<160], color="skyblue", label="Actual value")
sns.distplot(pred_test_y[pred_test_y<160], color="red", label="Predicted value")
plt.legend()

plt.tight_layout()
```



## k-fold Cross Validation using XGBoost

```
In [28]: dmatrix_train = xgb.DMatrix(data=train_data_feature_tf,label=train_data_target)

params = {'objective':'reg:linear', 'colsample_bytree': 0.3, 'learning_rate': 0.3, 'max_

model_cv = xgb.cv(dtrain=dmatrix_train, params=params, nfold=3, num_boost_round=50, earl
metrics="rmse", as_pandas=True, seed=7)

model_cv.tail(4)
```

```
[17:02:33] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-
030221e36e1a46bfb-1/xgboost/xgboost-ci-windows/src/objective/regression_obj.cu:213: reg:
linear is now deprecated in favor of reg:squarederror.
[17:02:33] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-
030221e36e1a46bfb-1/xgboost/xgboost-ci-windows/src/objective/regression_obj.cu:213: reg:
linear is now deprecated in favor of reg:squarederror.
[17:02:33] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-
030221e36e1a46bfb-1/xgboost/xgboost-ci-windows/src/objective/regression_obj.cu:213: reg:
linear is now deprecated in favor of reg:squarederror.
```

```
Out[28]:
```

	train-rmse-mean	train-rmse-std	test-rmse-mean	test-rmse-std
31	8.935207	0.183408	11.060047	0.736219
32	8.880285	0.174860	11.044372	0.740167
33	8.849045	0.185327	11.049080	0.738351
34	8.792400	0.202135	11.043289	0.728256

## Prediction on test data set using XGBoost

```
In [29]: #Assigning feature to test dataset
test_data_feature = test_data.drop(['ID'],axis=1)
print(test_data_feature.shape)
```

```
(4209, 364)
```

```
In [30]: #Describing object type variables in test dataset
test_data_feature.describe(include='object')
```

```
Out[30]:
```

	X0	X1	X2	X3	X4	X5	X6	X8
--	----	----	----	----	----	----	----	----

<b>count</b>	4209	4209	4209	4209	4209	4209	4209	4209
<b>unique</b>	49	27	45	7	4	32	12	25
<b>top</b>	ak	aa	as	c	d	v	g	e
<b>freq</b>	432	826	1658	1900	4203	246	1073	274

```
In [31]: test_data_feature['X0'] = le.fit_transform(test_data_feature.X0)
test_data_feature['X1'] = le.fit_transform(test_data_feature.X1)
test_data_feature['X2'] = le.fit_transform(test_data_feature.X2)
test_data_feature['X3'] = le.fit_transform(test_data_feature.X3)
test_data_feature['X4'] = le.fit_transform(test_data_feature.X4)
test_data_feature['X5'] = le.fit_transform(test_data_feature.X5)
test_data_feature['X6'] = le.fit_transform(test_data_feature.X6)
test_data_feature['X8'] = le.fit_transform(test_data_feature.X8)
```

```
In [32]: pca.fit(test_data_feature)
```

```
Out[32]: PCA(n_components=0.95)
```

```
In [33]: test_data_feature_tf = pca.fit_transform(test_data_feature)
print(test_data_feature_tf.shape)

(4209, 6)
```

```
In [34]: test_pred = model.predict(test_data_feature_tf)
test_pred
```

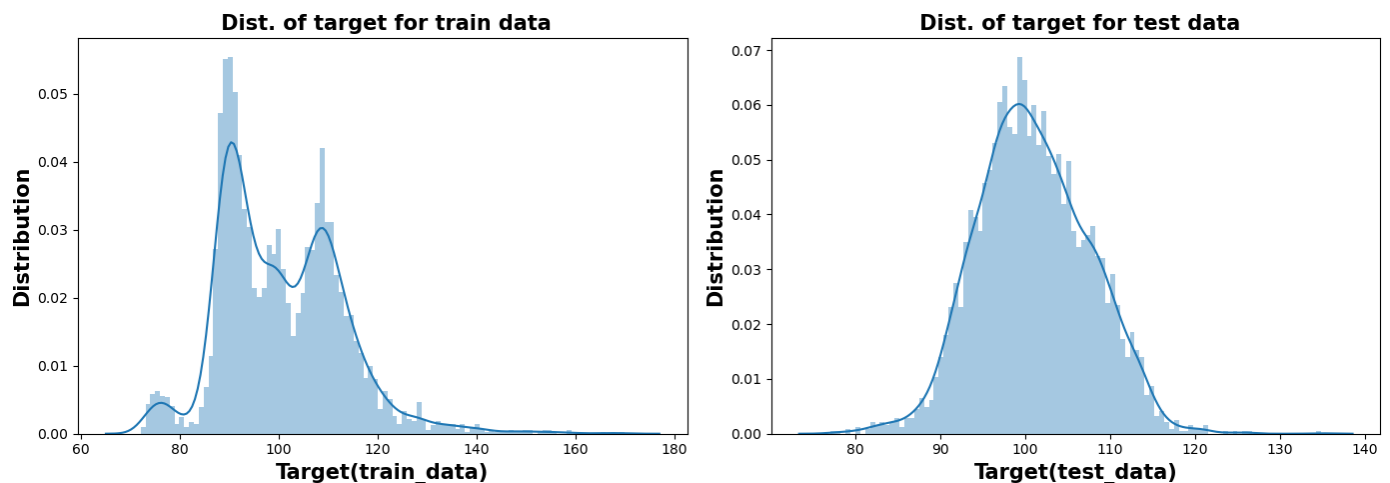
```
Out[34]: array([ 86.12015 ,  92.929794,  98.74635 , ...,  92.836525, 118.76457 ,
        98.46741 ], dtype=float32)
```

```
In [35]: fig, ax = plt.subplots(1,2, figsize=(14,5))

train_plot = sns.distplot(train_data_target[train_data_target<200], bins=100, kde=True,
train_plot.set_xlabel('Target(train_data)', weight='bold', size=15)
train_plot.set_ylabel('Distribution', weight='bold', size=15)
train_plot.set_title(' Dist. of target for train data', weight='bold', size=15)

test_plot = sns.distplot(test_pred[test_pred<200], bins=100, kde=True, ax=ax[1])
test_plot.set_xlabel('Target(test_data)', weight='bold', size=15)
test_plot.set_ylabel('Distribution', weight='bold', size=15)
test_plot.set_title(' Dist. of target for test data', weight='bold', size=15)

plt.tight_layout()
```



```
In [ ]:
```